What is Cosine Distance?

Cosine distance measures the dissimilarity between two vectors based on the cosine of the angle between them. It focuses purely on the direction of the vectors, ignoring their magnitude (length). The closer the vectors are in direction (smaller angle), the smaller the cosine distance (and higher the cosine similarity).

Think of it this way: if two vectors point in almost the same direction, they are considered very similar, even if one vector is much longer than the other.

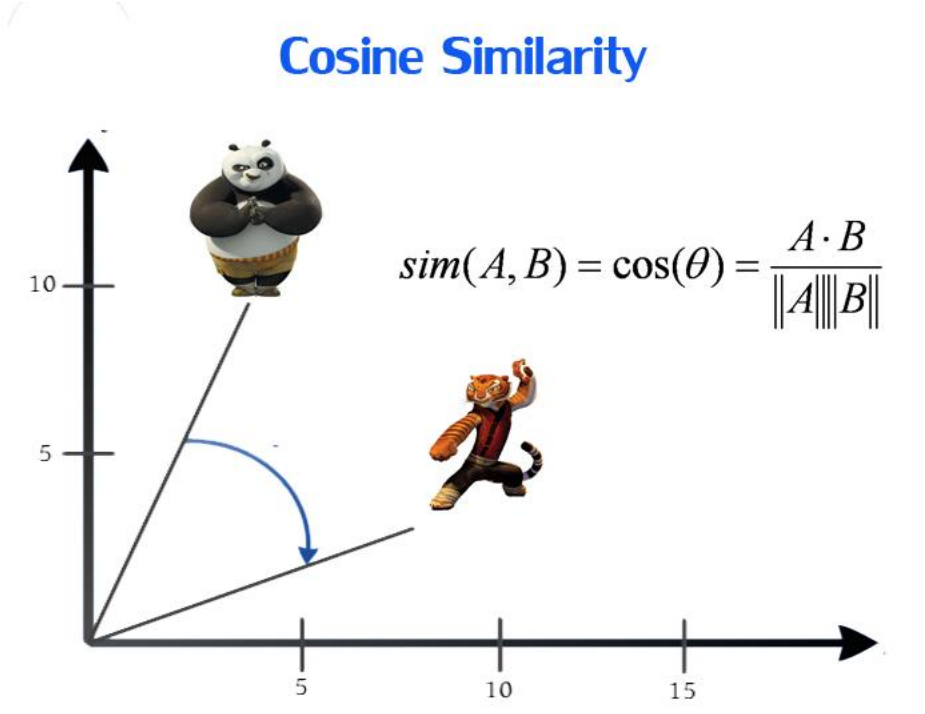

Formulas:

- Cosine Similarity: $similarity(A, B) = \frac{A \cdot B}{||A|| \cdot ||B||}$
  where $A \cdot B$ is the dot product of vectors A and B, and $||A||$ and $||B||$ are their magnitudes (lengths).
  - Ranges from -1 (completely opposite) to 1 (exactly the same direction). 0 means orthogonal (no relationship).
- Cosine Distance: $distance(A, B) = 1 - similarity(A, B)$
  - Ranges from 0 (exactly the same direction) to 2 (completely opposite). 1 means orthogonal.

Example: Customer Movie Preferences

Imagine a streaming service wants to recommend movies. They track how much a customer watches certain genres. Let's simplify and say we track two genres: Action and Comedy. The values represent hours watched in a month.

- Customer A: Watches 10 hours of Action, 2 hours of Comedy.
    - Vector A = (10, 2)
- Customer B: Watches 5 hours of Action, 1 hour of Comedy.
    - Vector B = (5, 1)
- Customer C: Watches 3 hours of Action, 9 hours of Comedy.
    - Vector C = (3, 9)

Let's calculate the Cosine Distance:

1. Customer A vs. Customer B:

- Dot Product ($A \cdot B$): $(10 \times 5) + (2 \times 1) = 50 + 2 = 52$
- Magnitude of A ($||A||$): $\sqrt{10^2 + 2^2} = \sqrt{100 + 4} = \sqrt{104} \approx 10.198$
- Magnitude of B ($||B||$): $\sqrt{5^2 + 1^2} = \sqrt{25 + 1} = \sqrt{26} \approx 5.099$
- Cosine Similarity (A, B): $\frac{52}{10.198 \times 5.099} \approx \frac{52}{51.99} \approx 1.00$ (very close to 1)
- Cosine Distance (A, B): $1 - 1.00 = 0.00$

Interpretation (A vs B): The cosine distance is very close to 0. This means Customer A and Customer B have very similar preferences in terms of the proportion of Action vs. Comedy they watch. Both watch 5 times more Action than Comedy. Even though Customer A watches more movies overall, their taste profile is identical.

2. Customer A vs. Customer C:

- Dot Product ($A \cdot C$): $(10 \times 3) + (2 \times 9) = 30 + 18 = 48$
- Magnitude of A ($||A||$): $\sqrt{104} \approx 10.198$ (already calculated)
- Magnitude of C ($||C||$): $\sqrt{3^2 + 9^2} = \sqrt{9 + 81} = \sqrt{90} \approx 9.487$
- Cosine Similarity (A, C): $\frac{48}{10.198 \times 9.487} \approx \frac{48}{96.75} \approx 0.496$
- Cosine Distance (A, C): $1 - 0.496 = 0.504$

Interpretation (A vs C): The cosine distance (0.504) is significantly larger than 0. This indicates that Customer A and Customer C have dissimilar preferences. Customer A prefers Action heavily, while Customer C prefers Comedy heavily. Their vectors point in very different directions.

When to Use Cosine Distance:

Cosine distance is ideal when:

- The magnitude (total quantity, like total hours watched, or total word count in a document) is not as important as the proportion or pattern of values across features.
- You are dealing with sparse data, such as in text analysis (where vectors represent word frequencies).
- You want to compare the "orientation" of user preferences, document topics, or image features.

Common Real-Life Applications

1. Information Retrieval

- **Document Similarity:** Cosine distance is used to compare documents and retrieve similar ones from a large corpus. For example, a search engine can use cosine distance to find documents that are semantically similar to a user's query.
- **Plagiarism Detection:** It is also used to detect plagiarism by comparing documents and flagging those with high cosine similarity.

2. Recommender Systems

- **Product Recommendation:** Cosine distance is used to recommend products to users based on their past purchases. For example, if two users have a high cosine similarity in their purchase history, the system may recommend products that one user has bought but the other hasn't.
- **Movie Recommendation:** It can also be used to recommend movies to users based on their ratings. For example, if two users have rated movies similarly, the system may recommend movies that one user has watched but the other hasn't.

3. Natural Language Processing (NLP)

- **Text Classification:** Cosine distance is used to classify text documents into different categories, such as spam or not spam, or news and sports.
- **Sentiment Analysis:** It is used to analyze the sentiment of a piece of text by comparing it to a known positive or negative sentiment. For example, a system can use cosine distance to determine if a review is positive or negative.
- **Word Embeddings:** In NLP, words are often represented as vectors in a high-dimensional space. Cosine distance is used to find words that are semantically similar to each other. For example, the words "king" and "queen" are likely to have a high cosine similarity.

Why it is Preferred

The reason cosine distance is preferred in these use cases is that it focuses on the **angle between the vectors**, not their magnitude. This is particularly useful for text data where the length of a document (and thus the magnitude of its vector) can vary greatly, but the topic or content remains the same.